

000
001
002
003
004
005
006
007
008
009
010
011
012
013
014
015
016
017
018
019
020
021
022
023
024
025
026
027
028
029
030
031
032
033
034
035
036
037
038
039
040
041
042
043
044
045
046
047
048
049
050
051
052
053
054

SWE-Edit: Rethinking Code Editing for Efficient SWE-Agent

Anonymous Authors¹

Abstract

Large language model agents have achieved remarkable progress on software engineering tasks, yet current approaches suffer from a fundamental *context coupling problem*: the standard code editing interface conflates code inspection, modification planning, and edit execution within a single context window, forcing agents to interleave exploratory viewing with strictly formatted edit generation. This causes irrelevant information to accumulate and degrades agent performance. To address this, we propose **SWE-Edit**, which decomposes code editing into two specialized sub-agents: a *Viewer* that extracts task-relevant code on demand, and an *Editor* that executes modifications from high-level plans—allowing the main agent to focus on reasoning while delegating context-intensive operations to clean context windows. We further investigate what makes an effective editing model: observing that the prevalent find-and-replace format is error-prone, we train Qwen3-8B with GRPO to adaptively select editing modes, yielding improved editing efficiency over single-format baselines. On SWE-bench Verified, SWE-Edit improves resolved rate by 2.1% while reducing inference cost by 17.9%. We additionally propose a code editing benchmark that reliably predicts downstream agentic performance, providing practical guidance for editing model selection.

1. Introduction

LLM-based coding agents can now solve real-world software engineering tasks by iteratively exploring codebases and refining solutions (Yang et al., 2024; Wang et al., 2024). Central to these systems is the *code editing interface* (Schluntz, 2025)—the tools through which agents

¹Anonymous Institution, Anonymous City, Anonymous Region, Anonymous Country. Correspondence to: Anonymous Author <anon.email@domain.com>.

Preliminary work. Under review by the International Conference on Machine Learning (ICML). Do not distribute.

inspect files and apply modifications. However, current interfaces suffer from a fundamental *context coupling problem*: they conflate code inspection, modification planning, and edit execution within a single context window, forcing agents to interleave exploratory viewing with strictly formatted edit generation.

This coupling creates a structural tension. For example, effective debugging requires broad exploration—viewing multiple files, tracing dependencies, testing hypotheses—yet each viewed file snippet persists in context regardless of its ultimate relevance. Meanwhile, generating correct code edits demands focused attention on precise locations and formats. Prior work establishes that LLM performance degrades when task-relevant information is buried within irrelevant context (Shi et al., 2023; Liu et al., 2024). For coding agents, exploration and precision are fundamentally at odds: a single agent cannot simultaneously optimize for comprehensive code understanding (which benefits from viewing many files) and reliable edit generation (which benefits from clean, focused context).

The code editing interface comprises two core operations that manifest this tension. The **view** operation allows agents to inspect file contents, but since agents cannot see code before viewing, they must explore incrementally—inevitably accumulating irrelevant context. The **edit** operation presents orthogonal challenges. The dominant *find-replace* format requires exact string matching; a single whitespace mismatch causes the edit to fail. The alternative *whole-file rewrite* avoids matching errors but incurs prohibitive token costs for long files. More fundamentally, reasoning about *what* to modify and generating *properly formatted* edit instructions are cognitively distinct capabilities: strong reasoning models like OpenAI o1 (Jaech et al., 2024) excel at describing solutions but frequently fail to produce correctly formatted edits (Gauthier, 2024a). Even GPT-5¹ exhibits notable formatting failure rates on the Aider Polyglot code editing benchmark (Gauthier, 2024b)—a reliability gap largely overlooked when evaluating agents on code editing tasks (Gauthier, 2024b; Jimenez et al., 2023).

We propose **SWE-Edit**, a framework that addresses context coupling by decomposing the code editing interface into spe-

¹Throughout this paper, GPT-5 refers to the model with reasoning effort set to high.

055
056
057
058
059
060
061
062
063
064
065
066
067
068
069
070
071
072
073
074
075
076
077
078
079
080
081
082
083
084
085
086
087
088
089
090
091
092
093
094
095
096
097
098
099
100
101
102
103
104
105
106
107
108
109

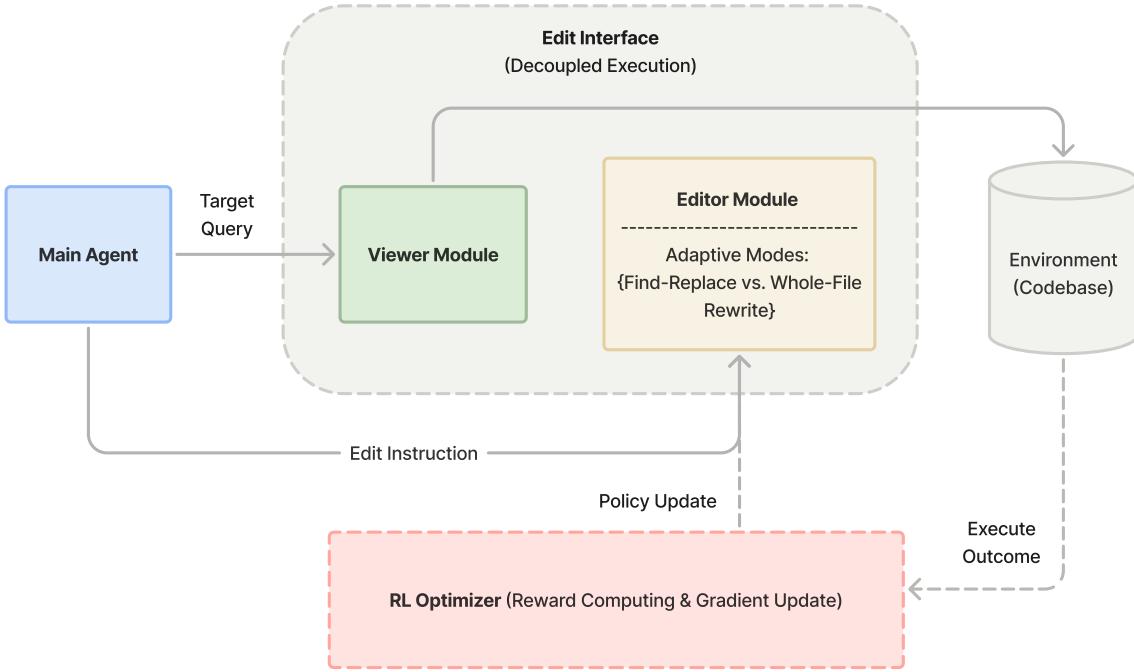


Figure 1. Overview of the proposed SWE-Edit framework architecture. The figure illustrates the **dual optimization** mechanism, demonstrating how optimization occurs simultaneously at both the *scaffolding level* (coordinating components and context) and the *model level* (refining the underlying models).

cialized subagents. A *Viewer* subagent receives complete files and extracts only task-relevant code on demand, eliminating exploratory context pollution from the main agent. An *Editor* subagent executes modifications from high-level natural language plans, decoupling reasoning from format-sensitive code generation. This decomposition allows the main agent to focus purely on problem-solving while delegating context-intensive operations to clean, specialized contexts.

Beyond scaffold design, we investigate what makes an effective editing model. Observing that the optimal editing strategy varies by find-replace suffices for small changes while whole-file rewrite handles complex restructuring—we train Qwen3-8B (Yang et al., 2025) with GRPO (Shao et al., 2024) to adaptively select editing modes based on modification complexity. The resulting model outperforms single-format baselines in editing accuracy. We further introduce a code editing benchmark and demonstrate that model performance on this benchmark reliably predicts downstream effectiveness when deployed as an editor subagent, providing practical guidance for editing model selection.

On SWE-bench Verified (Jimenez et al., 2023), our scaffolding-level contribution—the subagent decomposi-

tion—improves over the baseline by **2.1%** while reducing the inference cost by **17.9%**. By decoupling planning from format-sensitive generation, the decomposition also improves edit formatting reliability by **3.5%**. Our model-level contribution complements this by answering what makes an effective editor: the adaptive editing model and accompanying benchmark provide a principled path toward selecting and improving the editor subagent without costly end-to-end experimentation.

2. Related Work

LLM-Based Software Engineering Large language models have advanced from code completion (Chen, 2021; Austin et al., 2021; Jain et al., 2024; Zhuo et al., 2024) to code editing (Gauthier, 2024b) and repository-level software engineering (Jimenez et al., 2023). Early approaches to SWE tasks such as bug fixing employed fixed pipelines (Örwall, 2024; Xia et al., 2024), decomposing problems into localization, repair, and validation phases. Recent *agentic* systems (Yang et al., 2024; Wang et al., 2024) instead equip LLMs with tools for iterative codebase interaction. Our work advances this paradigm by redesigning the code editing interface—the central mechanism through

110 which SWE-agents inspect and modify code.

111
Multi-Agent Systems Multi-agent systems decompose
 112 complex problems across specialized agents. In software
 113 engineering, works like MetaGPT (Hong et al., 2023) and
 114 ChatDev (Qian et al., 2024) assign distinct development
 115 roles to communicating agents, while recent work (Hadfield
 116 et al., 2025) distributes research queries into independent,
 117 parallelizable subtasks with clear boundaries. These ap-
 118 proaches decompose at the *task or role* level—each agent
 119 pursues a *separable* objective. In contrast, SWE-Edit de-
 120 composes at the *cognitive* level: reasoning about *what* to
 121 modify and generating *properly formatted* edits are not inde-
 122 pendent subtasks but intertwined capabilities that interfere
 123 when sharing context. Our subagent design disentangles
 124 these conflicting cognitive demands within the code editing
 125 interface, and can integrate into broader multi-agent SWE
 126 frameworks.
 127

128
Training and Evolving SWE-Agents Parallel efforts im-
 129 prove SWE-agents through training or adaptation. Live-
 130 SWE-Agent (Xia et al., 2025) evolves scaffolding dynam-
 131 ically, while Context-Folding (Sun et al., 2025) trains
 132 agents to manage context via subagent delegation. SWE-
 133 Dev (Wang et al., 2025) and SWE-Fixer (Xie et al., 2025)
 134 advance open-source models through synthetic trajectories
 135 and retrieval-editing pipelines, respectively. In contrast,
 136 SWE-Edit requires no training and applies to closed-source
 137 models, though we also explore targeted training for the
 138 editing component rather than end-to-end agent behavior.
 139

140 3. Method: The SWE-Edit Framework

141
SWE-Edit adopts a two-stage optimization framework to
 142 improve the efficiency and reliability of code editing agents.
 143 At the *scaffolding level*, we decompose the code editing
 144 interface into specialized subagents, decoupling code in-
 145 spection from modification to reduce context pollution, and
 146 decoupling high-level reasoning from format-sensitive gen-
 147 eration to improve edit reliability (§3.1). At the *model*
 148 *level*, we train the editor to adaptively select between editing
 149 modes based on task characteristics, addressing the limita-
 150 tion that no single editing format performs optimally across
 151 modification types (§3.2).

152 3.1. Scaffolding Optimization

153 The standard code editing interface couples two distinct op-
 154 erations: inspecting code to understand context, and reason
 155 about modifying code to actually implement changes. This
 156 coupling forces the main agent to accumulate exploratory
 157 context that may be irrelevant to the final edit. SWE-Edit
 158 restructures the code editing interface by decomposing it
 159 into two subagents—**Viewer** and **Editor**—each operating
 160 in a clean, specialized context.

161
Viewer Subagent The viewer receives a file path and a nat-
 162 ural language query describing what information the agent
 163 seeks. Rather than returning raw file contents, it extracts
 164 and returns only the task-relevant code snippet. This filter-
 165 ing eliminates context pollution: the main agent receives
 166 precisely the information it needs without accumulating
 167 irrelevant code in its context window.

168
Editor Subagent The editor receives a file path and a nat-
 169 ural language edit instruction describing the desired mod-
 170 ification. It executes the edit directly, without requiring
 171 the main agent to produce format-sensitive find-replace
 172 commands. This decouples high-level reasoning—deciding
 173 *what* to change—from low-level generation—producing
 174 *correctly formatted* edit syntax.

175 Both subagents are implemented using a smaller, cost-
 176 efficient model, while the main agent focuses purely on
 177 problem-solving and orchestration. Full implementation
 178 details and prompts are provided in Appendix A.

179 3.2. Model Optimization

180 Given the scaffolding decomposition, a natural question
 181 arises: what makes an effective editor? As illustrated
 182 in Figure 2, the optimal editing strategy varies by task.
find-replace is token-efficient for localized changes but re-
 183 quires exact string matching—a single whitespace mismatch
 184 causes failure. *Whole-file rewrite* avoids matching errors
 185 but incurs higher cost and risks unintended modifications
 186 for long files. A static choice of either mode is suboptimal
 187 across heterogeneous editing tasks.

188 Algorithm 1 Lightweight Code Canonicalization

189 **Require:** Raw code string C

190 **Ensure:** Canonicalized code string \tilde{C}

191 1: **Step 1: Remove Comments**

192 2: Remove all multi-line comments from C

193 3: Remove all single-line comments from C

194 4: **Step 2: Normalize Whitespace**

195 5: Collapse whitespace in C to single space

196 6: $\tilde{C} \leftarrow \text{Trim}(C)$

197 7: **return** \tilde{C}

198 We address this by training the editor to *adaptively select*
 199 between editing modes. We formulate mode selec-
 200 tion as a single-step decision problem: given file con-
 201 tents c and edit instruction q , the editor chooses mode
 202 $m \in \{\text{find-replace}, \text{whole-file-rewrite}\}$ and
 203 generates the corresponding output. We optimize this policy
 204 using GRPO (Shao et al., 2024), with a *normalized match*

165
166
167
168
169
170
171
172
173
174
175
176
177
178
179
180
181
182
183
184

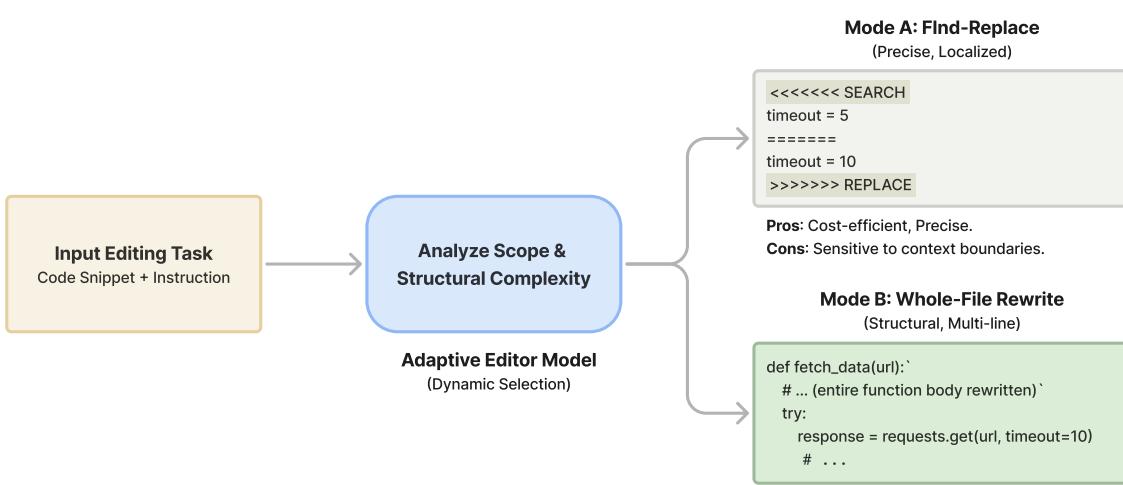


Figure 2. Adaptive editing mode selection. The editor analyzes task characteristics to choose between find-replace (token-efficient but matching-sensitive) and whole-file rewrite (robust but costly), enabling optimal strategy selection based on edit scope and complexity.

reward that compares model output against ground truth after canonicalizing whitespace and removing comments (Algorithm 1). This reward provides a reliable, execution-free proxy for edit correctness. Training details and the resulting benchmark are presented in §4.3.

4. Experiments

Our experiments are designed to validate the two core claims of SWE-Edit. We first examine **the effect of interface decomposition**, evaluating whether structurally decomposing the code editing interface into specialized subagents yields stable and consistent improvements in performance and cost efficiency over monolithic agent designs. We then study **model optimization under this decomposition**, investigating how a reinforcement learning-based adaptive editing strategy learns editing mode selection policies, thereby yielding more robust and efficient editor models.

4.1. General Setup

Evaluation We evaluate primarily on SWE-bench Verified (Jimenez et al., 2023), a curated benchmark of 500 real-world GitHub issues. All configurations are run 3 times to reduce variance. Full details of our SWE-Bench settings are provided in Appendix A.

Models For scaffolding experiments (§4.2), we use GPT-5 as the main agent and GPT-5-mini for both subagents by default. For model-level experiments (§4.3), we train

Qwen3-8B (Yang et al., 2025) with GRPO (Shao et al., 2024).

Metrics We report (1) **Resolve Rate**—percentage of issues where the agent’s patch passes all tests; (2) **Total Cost**—inference cost across all components; (3) **Edit Success Rate**—percentage of edit operations without formatting errors; (4) **Viewer/Editor Calls**—average number of viewer/editor tool calls per instance. We additionally track a detailed breakdown of token usage and inference cost. The full experiment table is presented in Appendix B.

4.2. Scaffolding-Level Results: Decomposition Yields Cost-Performance Synergy

A natural concern with subagent decomposition is the overhead of additional inference calls. We show that the viewer and editor subagents provide complementary benefits that combine synergistically: SWE-Edit achieves *both* higher resolve rate and lower cost than the monolithic baseline.

Viewer: Reducing Context Pollution We first evaluate the viewer subagent in isolation. As shown in Table 1, adding the viewer reduces total cost by 7.7% (\$243.7 → \$225.0) while slightly improving resolve rate (+0.4%). Two mechanisms drive this reduction: (1) the viewer uses a smaller model (GPT-5-mini) to process full file contents, and (2) by extracting only task-relevant snippets, it reduces context pollution in the main agent’s window. Notably, the average number of viewer calls decreases from 5.78 to 4.26 per instance; more robust responses streamline the process

220
221
222
223
224
225
226
227
228
229
230
231
Table 1. Main results on SWE-bench Verified (500 instances, 3 runs averaged). SWE-Edit improves resolve rate (+2.1%) and edit reliability (+3.5%) while reducing cost by 17.9%.

Configuration	Resolved (%)	Cost (\$)	Viewer Calls	Editor Calls	Edit Succ. (%)
Baseline	69.9	243.7	5.78	2.86	93.4
+ Viewer	70.3 (+0.4)	225.0 (-7.7%)	4.26 (-1.52)	2.75 (-0.11)	94.3 (+0.9)
+ Editor	71.3 (+1.4)	268.4 (+10.1%)	7.78 (+2.00)	2.33 (-0.53)	96.1 (+2.7)
SWE-Edit	72.0 (+2.1)	200.1 (-17.9%)	7.49 (+1.71)	2.37 (-0.49)	96.9 (+3.5)

230
231 and eliminate the need for back-and-forth file checking.

232
233 **Editor: Enhancing Edit Precision and Reliability** We
234 next evaluate the editor subagent in isolation. The inclusion
235 of the editor improves the SWE-bench Verified resolve rate
236 from 69.9% to 71.3% (**+1.4%**) and the edit success rate
237 from 93.4% to 96.1% (**+2.7%**). These results confirm that
238 decoupling high-level reasoning from syntactic execution
239 not only resolves formatting errors but also facilitates the
240 generation of logically sound patches. The average num-
241 ber of editor calls also decreases from 2.86 to 2.33 per
242 instance—higher edit reliability means fewer retry attempts
243 when edits fail due to incorrect patches. However, this reli-
244 ability gain comes at increased cost (**+10.1%**). Analyzing
245 agent trajectories, we find that the main agent becomes
246 more exploratory when delegating edits: viewer calls in-
247 crease from 5.78 to 7.78 per instance. We hypothesize that
248 generating natural language edit instructions, rather than
249 directly producing find-replace commands, encourages the
250 agent to gather more context before committing to an edit
251 plan. This behavioral shift further motivates combining the
252 editor with the viewer subagent.

253 **Synergistic Effect: Breaking the Accuracy-Cost Trade-**
254 **off** The complete SWE-Edit framework combines both
255 subagents, and the results reveal a synergistic effect: SWE-
256 Edit achieves the highest resolve rate (72.0%) at the lowest
257 cost (\$200.1), improving over baseline by **2.1%** absolute
258 while reducing cost by **17.9%**. Notably, SWE-Edit retains
259 the editor’s reliability benefits—edit success rate reaches
260 96.9% with only 2.37 editor calls per instance—while the
261 viewer offsets the increased exploration cost. Although
262 viewer calls remain elevated (7.49 vs. 5.78 baseline), the
263 viewer subagent processes these requests with a smaller
264 model and returns precise, task-relevant snippets rather than
265 raw file contents, yielding net cost savings despite more fre-
266 quent invocations. Figure 3 visualizes this trade-off: while
267 the viewer and editor individually improve one metric, some-
268 times at the expense of another, SWE-Edit uniquely occu-
269 pies the high-performance, low-cost quadrant.

270
271 **Generalization to Diverse Reasoning Models** To verify
272 that SWE-Edit’s benefits extend beyond GPT-5, we evaluate
273 on three recent reasoning models: Kimi-K2 (Moonshot AI,
274



275, MiniMax-M2.1 (Minimax AI, 2025), and GLM-
276 4.7 (ZhipuAI, 2025). Due to computational constraints,
277 we run each configuration twice on the first 100 instances
278 of SWE-bench Verified. Detailed settings are provided in
279 Appendix D.

280
281 Table 2. Generalization results on different reasoning models (100
282 instances, 2 runs). SWE-Edit consistently improves resolve rate
283 and substantially increases edit success rate across all models.

Model	Config	Resolved (%)	Edit Succ. (%)
Kimi K2 Thinking	Baseline	56.7	75.6
	SWE-Edit	59.4 (+2.7)	93.9 (+18.3)
MiniMax-M2.1	Baseline	58.8	82.0
	SWE-Edit	62.9 (+4.1)	94.8 (+12.8)
GLM-4.7	Baseline	63.3	79.6
	SWE-Edit	64.9 (+1.6)	95.9 (+16.3)

284 As shown in Table 2, SWE-Edit consistently outperforms the
285 baseline across all three models. The most striking improve-
286 ment is in edit success rate: while baseline configurations

275 *Table 3.* Results on PR-Edit Benchmark. GRPO training substantially improves Qwen3-8B, achieving performance comparable to
 276 GPT-5-nano.
 277

Model	Format (%)	GPT Grader (%)	Norm. Match (%)
Qwen3-8B	76.8	56.0	32.0
Qwen3-8B + GRPO	90.4	68.4	38.8
GPT-5-nano	89.8	66.4	38.8
GPT-5-mini	96.1	77.5	41.7
GPT-5	98.1	77.2	44.1

285 *Table 4.* Downstream performance on SWE-bench Verified with different editor models. Higher PR-Edit scores predict better resolve rate,
 286 higher edit success, and lower main agent cost.
 287

Editor Model	PR-Edit (%)	Resolved (%)	Agent Cost (\$)	Edit Succ. (%)
Qwen3-8B	56.0	68.5	231.7	68.6
Qwen3-8B + GRPO	68.4	69.9	215.9	81.1
GPT-5-nano	66.4	70.0	207.1	82.0
GPT-5-mini	77.5	72.0	179.6	95.9

295 exhibit variable reliability (75.6%–82.0%), SWE-Edit stabilizes performance at 93.9%–95.9%, representing gains of
 296 12.8–18.3 percentage points. This confirms that the editor
 297 subagent’s reliability benefits are model-agnostic reasoning
 298 models, which often struggle more with strict formatting
 299 requirements, benefit even more from the decoupled archi-
 300 tecture. Resolve rate improvements range from +1.6% to
 301 +4.1%, demonstrating that the scaffolding-level gains ob-
 302 served with GPT-5 transfer to other alternatives.
 303

4.3. Model-Level Results: Training Adaptive Editors

305 Having established that SWE-Edit yields robust benefits
 306 across model families, we now turn to the optimization
 307 of the editor itself. Under the proposed decomposition,
 308 a natural question arises: **How can we effectively train**
 309 and select models to excel in the specialized role of an
 310 editor? To address this, we leverage the modularity of
 311 our framework to perform targeted reinforcement learning
 312 on a Qwen3-8B-based backbone, transforming it into an
 313 adaptive, high-precision editing subagent. We compare the
 314 resulting model against the original Qwen3-8B baseline and
 315 evaluate both on our PR-Edit benchmark and SWE-bench
 316 Verified.
 317

4.3.1. TRAINING SETUP

318 **Data** We curate training data from open-source GitHub
 319 pull requests across diverse repositories. For each PR, we
 320 extract the file content before and after merging, along with
 321 the git diff and PR message. We prompt GPT-4.1 to gen-
 322 erate natural language edit instructions conditioned on these
 323 artifacts, yielding 3.5K examples split into 2.8K training,
 324 200 validation, and 500 held-out test instances.
 325

RL Training We fine-tune Qwen3-8B (Yang et al., 2025) using GRPO (Shao et al., 2024) on Slime (Zhu et al., 2025). The model learns to select an editing mode and generate the corresponding output. Training runs for 520 rollout steps; each step samples 32 instances with 8 candidate outputs per instance. We use the normalized match reward described in §3.2.

Intermediate Evaluation (PR-Edit Benchmark) We re-serve the 500 held-out examples as the **PR-Edit Benchmark**, which serves as a lightweight and efficient intermediate evaluation for editor models. Compared to end-to-end evaluation on SWE-bench Verified, PR-Edit is substantially cheaper and faster to run, enabling rapid iteration and controlled analysis of editor behavior under the proposed decomposition. Importantly, performance on PR-Edit is strongly correlated with downstream agent performance on SWE-bench Verified, as evidenced by the results reported in Tables 3 and 4. We report three metrics: (1) *Format Success*, measuring whether the output conforms to the required editing format; (2) *GPT Grader*, where GPT-4.1 assesses whether the generated edit correctly implements the instruction; and (3) *Normalized Match*, computing exact match after canonicalization.

4.3.2. EVALUATION OF ADAPTIVE EDITOR TRAINING

We evaluate the effectiveness of adaptive editor training at both the editor level and in downstream agent performance. Table 3 reports results on the PR-Edit Benchmark, which provides a controlled evaluation of editor behavior under the proposed decomposition. GRPO training substantially improves Qwen3-8B as an editor model: format success increases from 76.8% to 90.4% (**+13.6%**), and GPT Grader

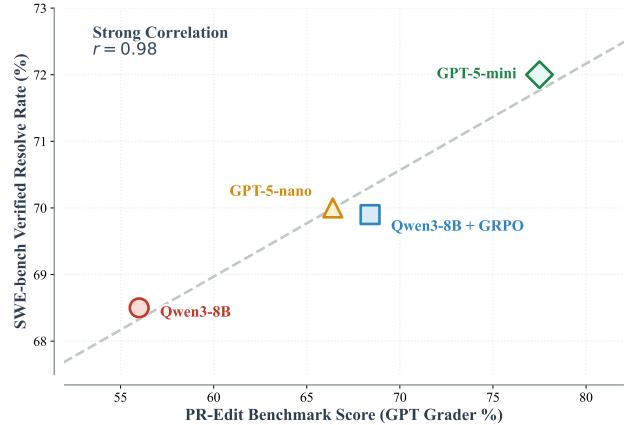


Figure 4. PR-Edit benchmark scores correlate with downstream agent performance, enabling efficient editor model selection without full SWE-bench evaluation.

accuracy improves from 56.0% to 68.4% (**+12.4%**). Across all reported metrics, the trained model exceeds GPT-5-nano.

We next examine whether these editor-level improvements translate into gains when the trained model is deployed as an editor subagent within SWE-Edit. This question is practically important, as end-to-end evaluation on SWE-bench Verified is costly. As shown in Table 4, improvements on the PR-Edit Benchmark consistently correspond to stronger downstream performance, including higher resolve rates, higher edit success rates, and lower main-agent inference cost. In particular, the GRPO-trained Qwen3-8B improves the SWE-bench Verified resolve rate from 69.9% to 71.3% (**+1.4%**), while reducing main-agent inference cost by 6.8%. These gains are driven by a substantial increase in edit success rate from 68.6% to 81.1% (**+12.5%**), indicating that editor-level improvements reliably translate into end-to-end agent effectiveness.

Finally, Figure 4 visualizes the relationship between PR-Edit performance and downstream SWE-bench effectiveness, demonstrating a strong correlation between editor-level benchmark scores and end-to-end agent performance. Together, these results show that adaptive editor training yields consistent improvements and that PR-Edit serves as a reliable intermediate signal for editor quality.

Together, these results show that adaptive editor training yields consistent improvements and that PR-Edit serves as a reliable intermediate signal for editor quality.

4.4. Ablation Studies

Adaptive vs. Fixed Edit Format We compare our adaptive format selection against a fixed find-replace baseline. We focus on find-replace as the fixed baseline because it is the dominant choice in practice—used by Claude’s

Table 5. Ablation on edit format strategy. Adaptive selection outperforms fixed find-replace by learning to match strategy to task complexity. All metrics reported in percentages except Cost (in USD).

Strategy	PR-Edit	Resolved	Cost	Edit Succ.
Search-Replace	67.0	69.4	244.7	80.2
Adaptive (Ours)	68.4	69.9	215.9	81.1

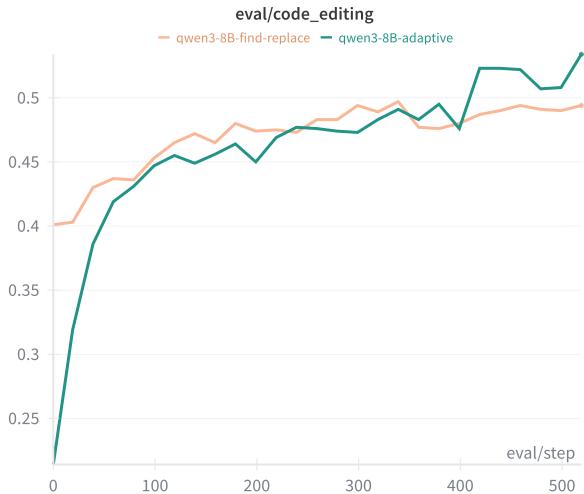


Figure 5. Training dynamics for fixed vs. adaptive format selection. While fixed find-replace starts higher (simpler format, easier to learn), adaptive training surpasses it by learning when to invoke whole-file rewrite.

`str_replace_editor` (Schlutz, 2025), Aider (Gauthier, 2024b), and most agentic coding systems. The alternative, whole-file rewrite, is impractical as a sole strategy: real-world code files often span thousands of lines, making full rewrites prohibitively expensive and prone to introducing unintended changes.

Table 5 shows that adaptive selection outperforms the fixed baseline on both PR-Edit benchmark and downstream SWE-bench evaluation. The improvement stems from task heterogeneity: localized bug fixes and single-line changes are handled efficiently by find-replace, while complex refactoring or multi-site edits benefit from whole-file rewrite where precise string matching is error-prone.

Figure 5 illustrates the training dynamics. The fixed find-replace policy starts with higher validation reward—unsurprising, as find-replace is simpler and most training examples involve localized edits. However, adaptive training converges to a higher final reward, as the model learns to invoke whole-file rewrite for the subset of tasks where find-replace struggles. This confirms our hypothesis: a single format cannot optimally serve all edit types, and learning to select adaptively yields consistent gains.

385
386
387
388
389
390
391
392
393
394
395
396
397
398
399
400
401
402
403
404
405
406
407
408
409
410
411
412
413
414
415
416
417
418
419
420
421
422
423
424
425
426
427
428
429
430
431
432
433
434
435
436
437
438
439

Table 6. Effect of editor model scale. Stronger models show diminishing returns: GPT-5 provides minimal accuracy gain at $5.8 \times$ the cost. Percentages (%) and cost (\$) omitted from headers for brevity.

Editor Model	Resolved	Edit Succ.	Editor Cost
GPT-5-mini	72.0	95.9	5.4
GPT-5	72.4 (+0.4)	97.5 (+1.6)	31.2 ($5.8 \times$)

Scaling the Editor Model We examine whether stronger models yield proportional gains in the editor role by replacing GPT-5-mini with GPT-5. As shown in Table 6, GPT-5 improves resolve rate by only 0.4% (72.0% → 72.4%) while increasing editor cost by $5.8 \times$ ($\$5.4 \rightarrow \31.2 per run). This diminishing return suggests that the editor role is bottlenecked by format reliability rather than reasoning capability: once edit success rate saturates near 98%, additional model capacity provides marginal benefit. For cost-sensitive deployments, smaller models suffice for the editor role, while investment in the main agent yields higher returns.

4.5. Summary of Findings

Taken together, our experimental results validate both core claims of SWE-Edit: that interface decomposition yields stable and consistent gains in performance and cost efficiency, and that reinforcement learning-based adaptive editor optimization enables the learning of robust and efficient editing behaviors under this decomposition.

At the **scaffolding level**, experimental results in Table 1 show that decomposing the code editing interface into functionally specialized subagents yields consistent performance improvements while substantially reducing inference cost. Specifically, this decomposition improves the overall resolve rate by 2.1% and reduces inference cost by 17.9%. Further analysis reveals that the Viewer and Editor subagents provide complementary and non-substitutable benefits. The Viewer mitigates context pollution by selectively extracting decision-relevant code snippets on demand, while the Editor improves execution stability and success rate by decoupling high-level reasoning from the format-sensitive edit generation process. Their combination yields structural gains that cannot be achieved by either component alone, validating the effectiveness of the proposed interface decomposition at the system level.

At the **model level**, results in Tables 3 and 4 further demonstrate that performance on the PR-Edit benchmark reliably predicts downstream agent effectiveness on SWE-bench. PR-Edit thus serves as an efficient and stable intermediate proxy for guiding editor model selection and optimization. Compared to the untrained Qwen3-8B baseline, models trained with adaptive editor optimization (GRPO) achieve significant and consistent improvements at both the editor

level and in downstream SWE-bench performance. These gains indicate that downstream improvements primarily stem from adaptive optimization of editing behavior, rather than from increases in model scale or inference budget. Overall, these results highlight the critical role of adaptive editor training in improving both editing reliability and end-to-end agent performance, and point to a practical path toward building efficient and scalable code editing subagents.

Finally, experiments on reasoning models show consistent improvements across model families in both resolve rate and edit reliability.

5. Conclusions and Future Work

We present SWE-Edit, a two-stage optimization framework to enhance the code editing interface in software engineering agents. Our scaffolding-level decomposition improves agents’ performance in SWE tasks while saving significant inference cost. At the model level, we showed that training Qwen3-8B with GRPO to adaptively select editing modes yields substantial improvements over single-format baselines, and introduced the PR-Edit benchmark as a reliable proxy for downstream agent effectiveness.

One limitation of our current approach is that we train the editor model in isolation using a static reward signal derived from ground-truth edits. A natural extension is to train the editor within an end-to-end agentic reinforcement learning loop, where it receives feedback from the main agent’s downstream success or failure. This would allow the editor to learn not just formatting correctness, but also properties that facilitate effective agent-editor collaboration—such as producing edits that are easier for the main agent to verify or debug when errors occur. More broadly, our subagent decomposition provides a modular foundation for investigating how specialized components should co-evolve within multi-agent software engineering systems.

Impact Statement

This paper presents work whose goal is to advance the field of Machine Learning. There are many potential societal consequences of our work, none which we feel must be specifically highlighted here.

References

- Austin, J., Odena, A., Nye, M., Bosma, M., Michalewski, H., Dohan, D., Jiang, E., Cai, C., Terry, M., Le, Q., et al. Program synthesis with large language models. *arXiv preprint arXiv:2108.07732*, 2021.

- Chen, M. Evaluating large language models trained on code. *arXiv preprint arXiv:2107.03374*, 2021.

- 440 Gauthier, P. Aider architect: Separating code reasoning
 441 and editing. <https://aider.chat/2024/09/26/architect.html>, September 2024a.
- 442
- 443 Gauthier, P. Aider polyglot benchmark. <https://aider.chat/docs/leaderboards/>, 2024b.
- 444
- 445 Hadfield, J., Zhang, B., Lien, K., Scholz,
 446 F., Fox, J., and Ford, D. How we built
 447 our multi-agent research system. <https://www.anthropic.com/engineering/multi-agent-research-system>, June 2025.
- 448
- 449 Hong, S., Zhuge, M., Chen, J., Zheng, X., Cheng, Y., Wang,
 450 J., Zhang, C., Wang, Z., Yau, S. K. S., Lin, Z., et al.
 451 Metagpt: Meta programming for a multi-agent collaborative
 452 framework. In *The twelfth international conference on learning representations*, 2023.
- 453
- 454 Jaech, A., Kalai, A., Lerer, A., Richardson, A., El-Kishky,
 455 A., Low, A., Helyar, A., Madry, A., Beutel, A., Carney,
 456 A., et al. Openai o1 system card. *arXiv preprint arXiv:2412.16720*, 2024.
- 457
- 458 Jain, N., Han, K., Gu, A., Li, W.-D., Yan, F., Zhang, T.,
 459 Wang, S., Solar-Lezama, A., Sen, K., and Stoica, I.
 460 Livecodebench: Holistic and contamination free eval-
 461 uation of large language models for code. *arXiv preprint arXiv:2403.07974*, 2024.
- 462
- 463 Jimenez, C. E., Yang, J., Wettig, A., Yao, S., Pei, K., Press,
 464 O., and Narasimhan, K. Swe-bench: Can language mod-
 465 els resolve real-world github issues? *arXiv preprint arXiv:2310.06770*, 2023.
- 466
- 467 Liu, N. F., Lin, K., Hewitt, J., Paranjape, A., Bevilacqua,
 468 M., Petroni, F., and Liang, P. Lost in the middle: How
 469 language models use long contexts. *Transactions of the Association for Computational Linguistics*, 12:157–173, 2024.
- 470
- 471 MiniMax AI. MiniMax M2.1: Significantly enhanced
 472 multi-language programming, built for real-world com-
 473 plex tasks. <https://www.minimax.io/news/minimax-m21>, December 2025.
- 474
- 475 Moonshot AI. Kimi k2 thinking. <https://moonshotai.github.io/Kimi-K2/thinking.html>, 2025.
- 476
- 477 Qian, C., Liu, W., Liu, H., Chen, N., Dang, Y., Li, J., Yang,
 478 C., Chen, W., Su, Y., Cong, X., et al. Chatdev: Commu-
 479 nicative agents for software development. In *Proceedings of the 62nd Annual Meeting of the Association for Computational Linguistics (Volume 1: Long Papers)*, pp. 15174–15186, 2024.
- 480
- 481 Schluntz, E. Raising the bar on swe-bench verified with
 482 claude 3.5 sonnet. Anthropic Engineering Blog, Jan-
 483 uary 2025. URL <https://www.anthropic.com/engineering/swe-bench-sonnet>.
- 484
- 485 Shao, Z., Wang, P., Zhu, Q., Xu, R., Song, J., Bi, X., Zhang,
 486 H., Zhang, M., Li, Y., Wu, Y., et al. Deepseekmath: Push-
 487 ing the limits of mathematical reasoning in open language
 488 models. *arXiv preprint arXiv:2402.03300*, 2024.
- 489
- 490 Shi, F., Chen, X., Misra, K., Scales, N., Dohan, D., Chi,
 491 E. H., Schärli, N., and Zhou, D. Large language models
 492 can be easily distracted by irrelevant context. In *Inter-
 493 national Conference on Machine Learning*, pp. 31210–
 494 31227. PMLR, 2023.
- 495
- 496 Sun, W., Lu, M., Ling, Z., Liu, K., Yao, X., Yang, Y., and
 497 Chen, J. Scaling long-horizon llm agent via context-
 498 folding. *arXiv preprint arXiv:2510.11967*, 2025.
- 499
- 500 Wang, H., Hou, Z., Wei, Y., Tang, J., and Dong, Y. Swe-
 501 dev: Building software engineering agents with training
 502 and inference scaling. *arXiv preprint arXiv:2506.07636*,
 503 2025.
- 504
- 505 Wang, X., Li, B., Song, Y., Xu, F. F., Tang, X., Zhuge, M.,
 506 Pan, J., Song, Y., Li, B., Singh, J., et al. Openhands: An
 507 open platform for ai software developers as generalist
 508 agents. *arXiv preprint arXiv:2407.16741*, 2024.
- 509
- 510 Xia, C. S., Deng, Y., Dunn, S., and Zhang, L. Agentless: De-
 511 mystifying llm-based software engineering agents. *arXiv preprint arXiv:2407.01489*, 2024.
- 512
- 513 Xia, C. S., Wang, Z., Yang, Y., Wei, Y., and Zhang, L. Live-
 514 swe-agent: Can software engineering agents self-evolve
 515 on the fly? *arXiv preprint arXiv:2511.13646*, 2025.
- 516
- 517 Xie, C., Li, B., Gao, C., Du, H., Lam, W., Zou, D., and
 518 Chen, K. Swe-fixer: Training open-source llms for effec-
 519 tive and efficient github issue resolution. *arXiv preprint arXiv:2501.05040*, 2025.
- 520
- 521 Yang, A., Li, A., Yang, B., Zhang, B., Hui, B., Zheng, B.,
 522 Yu, B., Gao, C., Huang, C., Lv, C., et al. Qwen3 technical
 523 report. *arXiv preprint arXiv:2505.09388*, 2025.
- 524
- 525 Yang, J., Jimenez, C. E., Wettig, A., Lieret, K., Yao, S.,
 526 Narasimhan, K., and Press, O. Swe-agent: Agent-
 527 computer interfaces enable automated software engineer-
 528 ing. *Advances in Neural Information Processing Systems*,
 529 37:50528–50652, 2024.
- 530
- 531 ZhipuAI. GLM-4.7: Advancing the coding capability.
 532 <https://z.ai/blog/glm-4.7>, December 2025.
- 533
- 534 Zhu, Z., Xie, C., Lv, X., and slime Contributors. slime: An
 535 llm post-training framework for rl scaling. <https://>

495 github.com/THUDM/slime, 2025. GitHub repository. Corresponding author: Xin Lv.

496
497 Zhuo, T. Y., Vu, M. C., Chim, J., Hu, H., Yu, W., Widyasari,
498 R., Yusuf, I. N. B., Zhan, H., He, J., Paul, I., et al. Big-
499 codebench: Benchmarking code generation with diverse
500 function calls and complex instructions. *arXiv preprint*
501 *arXiv:2406.15877*, 2024.

502
503 Örwall, A. Moatless. <https://github.com/aorwall/moatless-tools>, 2024.

504

505

506

507

508

509

510

511

512

513

514

515

516

517

518

519

520

521

522

523

524

525

526

527

528

529

530

531

532

533

534

535

536

537

538

539

540

541

542

543

544

545

546

547

548

549

550 A. Implementation Details

551 This appendix provides full details of the agent scaffolding, tool definitions, and prompts used in our experiments.

553 554 A.1. Baseline Agent Scaffolding

555 We adopt the reference agent scaffolding from Anthropic ([Schluntz, 2025](#)), which equips the agent with two tools:
 556 `execute_bash` for shell command execution and `str_replace_editor` for file operations. The editor tool provides
 557 sub-commands for viewing, creating, and editing files via exact string replacements. This reflects current best practices for
 558 agentic software engineering.
 559

560 **Tool Definitions.** The baseline agent uses the following tools:

561 *Listing 1.* Bash tool schema.

```
562
563 {
564     "type": "function",
565     "name": "execute_bash",
566     "description": "Run commands in a bash shell\n* When invoking this tool, the contents of the \"command\" parameter  

567         does NOT need to be XML-escaped.\n* You don't have access to the internet via this tool.\n* You do have  

568         access to a mirror of common linux and python packages via apt and pip.\n### Command Execution\n* **Non-  

569         persistent**: Each shell tool call is executed in a fresh environment. Shell variables, working directory  

570         changes, and history are NOT preserved between calls.\n* **Timeout**: Commands have a default timeout of 120  

571         seconds (max 300). Set the 'timeout' parameter for long-running commands.\n* **One command at a time**: Chain  

572         multiple commands using `&&` (conditional), `;` (sequential), or `||` (on failure).\n### Long-running  

573         Commands\n* For commands that may run indefinitely (e.g., servers), run in background: 'python3 app.py >  

574         server.log 2>& 1 &\n* For potentially long commands (installations, tests), set an appropriate 'timeout'  

575         value.\n### Best Practices\n* **Avoid large outputs**: Commands producing massive output may be truncated.\n* **Directory verification**: Verify parent directories exist before creating/editing files.\n### Output  

576         Handling\n* Stdout and stderr are combined and returned as a string. Output may be truncated if too long.\n*  

577         Exit codes are provided in system tags for failed commands.\n* Timeout messages are returned if commands  

578         exceed the timeout limit.",
579     "parameters": {
580         "type": "object",
581         "properties": {
582             "command": {
583                 "type": "string",
584                 "description": "The bash command to execute. You can only execute one bash command at a time. If you need to  

585                     run multiple commands sequentially, you can use `&&` or `;` to chain them together."
586             },
587             "timeout": {
588                 "type": "integer",
589                 "description": "Optional timeout in seconds for the command execution. If the command takes longer than this  

590                     , it will be terminated.",
591                 "default": 120,
592                 "minimum": 1,
593                 "maximum": 300
594             }
595         },
596         "required": ["command"]
597     }
598 }
```

599 The `str_replace_editor` tool in the baseline follows the standard Anthropic implementation for file viewing and
 600 editing via string matching. We omit its full schema for brevity as it follows the standard specification in [Schluntz \(2025\)](#).

601 602 A.2. SWE-Edit Agent Scaffolding

603 SWE-Edit replaces the `str_replace_editor` with a unified `llm_editor` tool that internally delegates to viewer and
 604 editor subagents. The `execute_bash` tool remains unchanged.

605 **Tool Definition.** The `llm_editor` tool exposes a unified interface for file viewing, creation, and editing, with AI-powered
 606 processing for the `view` and `edit` commands.

607 *Listing 2.* LLM editor tool schema.

```
608
609 {
610     "type": "function",
611     "name": "llm_editor",
612     "description": "Custom editing tool for viewing, creating and editing files\n* State is persistent across command  

613         calls and discussions with the user\n* If 'path' is a directory, 'view' lists non-hidden files and
```

```

605     directories up to 2 levels deep\n* If 'path' is a file, 'view' uses AI to find and display only the sections
606     relevant to your 'query'\n* The 'create' command cannot be used if the specified 'path' already exists as a
607     file\n\nNotes for using the 'view' command:\n* Provide a 'query' describing what you're looking for (e.g., \"\n* Where is user authentication handled?\", \"Show me the class definition for User\")\n* The tool reads the
608     file and uses AI to identify relevant line ranges, then displays those sections with line numbers\n* Multiple
609     relevant sections are shown with '... (N lines omitted) ...' separators between them\n\nNotes for using the
610     'edit' command:\n* Provide a clear 'instruction' describing what to change and where (identify by function/
611     class/method name)\n* The tool reads the file internally and applies your instruction using AI-powered search
612     -replace\n* Be specific: \"In 'MyClass.my_method', change X to Y\" is better than \"fix the bug\"\n* After
613     editing, the output shows the modified regions",
614     "parameters": {
615       "type": "object",
616       "properties": {
617         "command": {
618           "type": "string",
619           "enum": ["view", "create", "edit"],
620           "description": "The command to run. Allowed options are: 'view', 'create', 'edit'."
621         },
622         "path": {
623           "type": "string",
624           "description": "Absolute path to file or directory, e.g. '/workspace/file.py' or '/workspace'."
625         },
626         "query": {
627           "type": ["string", "null"],
628           "default": null,
629           "description": "Required for 'view' command when 'path' points to a file. A natural language query
630               describing what you're looking for in the file. An LLM will analyze the file and return only the line
631               ranges relevant to your query. Examples: 'Where is the authentication logic?', 'Show me the class
632               definition for User', 'Find all functions that handle HTTP requests'."
633         },
634         "instruction": {
635           "type": ["string", "null"],
636           "default": null,
637           "description": "Required for 'edit' command. Detailed instruction describing how to modify the file. Be
638               specific about what changes to make and where (function/class/method name).."
639         },
640         "file_text": {
641           "type": ["string", "null"],
642           "default": null,
643           "description": "Required for 'create' command. The content of the file to be created."
644         }
645       },
646       "required": ["command", "path"]
647     }
648   }
649 
```

637 **Viewer Subagent Prompt.** When the `view` command is invoked on a file, the viewer subagent receives the file contents
638 with line numbers and the user's query. It returns relevant line ranges as a JSON array. The complete system prompt is
639 shown below.

640 **Viewer Subagent System Prompt**

641 You are an expert code analyzer. Your task is to identify line ranges
642 in a file that are relevant to a given query.

643 You will be given:

- 644 1. A file with numbered lines in the format: LINE_NUMBER\tLINE_CONTENT
645 2. A query describing what the user is looking for

646 Your job is to analyze the file and return the line ranges that are
647 most relevant to the query. Consider:

- 648 - Function/method definitions that match the query
- 649 - Class definitions related to the query
- 650 - Variable declarations or assignments relevant to the query
- 651 - Import statements if they're relevant
- 652 - Comments that explain relevant code
- 653 - Any code blocks that implement functionality related to the query

654 **OUTPUT FORMAT:**

655 You must output your response as a JSON array of line ranges. Each
656 range is an array of two integers [start_line, end_line] (inclusive,

```

660
661     1-indexed).
662
663     Example output:
664     [[10, 25], [45, 60], [100, 115]]
665
666     RULES:
667     1. Only output the JSON array, no additional explanation or comments
668     2. Line numbers are 1-indexed (first line is line 1)
669     3. Each range should include complete logical blocks (don't cut
670         functions/classes in the middle)
671     4. Include a few lines of context before and after each relevant
672         section when appropriate
673     5. If nothing in the file is relevant to the query, return an
674         empty array: []
675     6. Ranges should be sorted by start line number
676     7. Merge overlapping or adjacent ranges
677     8. Keep ranges focused - don't include entire files unless the
678         query asks for everything
679
680     Example 1 - Finding a specific function:
681     Query: "Where is the calculate_total function defined?"
682     Output: [[15, 28]]
683
684     Example 2 - Finding multiple related sections:
685     Query: "How is user authentication handled?"
686     Output: [[5, 8], [23, 45], [102, 130]]
687
688     Example 3 - Nothing relevant found:
689     Query: "Where is the database connection configured?"
690     Output: []
691
692     Now, analyze the file content and query provided, and output the
693     relevant line ranges as a JSON array.

```

690 **Editor Subagent Prompt.** When the edit command is invoked, the editor subagent receives the file contents and the
 691 edit instruction. It outputs modifications in search-replace format or rewrites the entire file when appropriate. The complete
 692 system prompt is shown below.
 693

694 Editor Subagent System Prompt

695 You are an expert code editor. Your task is to analyze a file and
 696 make modifications according to the provided instructions.
 697

698 You must output your changes using the search-replace format shown
 699 below. You can make multiple edits by including multiple
 700 search-replace blocks.
 701

702 Format for each edit:
 703 <<<<< SEARCH
 704 exact lines from the original file to find
 705 =====
 706 new lines to replace them with
 707 >>>>> REPLACE

708 **IMPORTANT RULES:**
 709 1. The SEARCH block must match the original file content EXACTLY,
 including whitespace and indentation
 710 2. You can make multiple edits by including multiple search-replace
 blocks
 711 3. If the SEARCH block is empty (no content between <<<<< SEARCH
 712 and =====), it means you want to REWRITE THE ENTIRE FILE with
 713

```

715     the content in the REPLACE block
716 4. Each SEARCH block must be unique in the file - if there are
717     multiple matches, include more context
718 5. Only output the search-replace blocks, no additional explanation
719     or comments
720
721 Example 1 - Modifying specific lines:
722 <<<<< SEARCH
723 def calculate_total(items):
724     return sum(items)
725 =====
726 def calculate_total(items):
727     if not items:
728         return 0
729     return sum(items)
730 >>>>> REPLACE
731
732 Example 2 - Multiple edits:
733 <<<<< SEARCH
734 import os
735 =====
736 import os
737 import sys
738 >>>>> REPLACE
739
740 <<<<< SEARCH
741 def main():
742     pass
743 =====
744 def main():
745     print("Hello, World!")
746 >>>>> REPLACE
747
748 Example 3 - Rewriting entire file (empty SEARCH block):
749 <<<<< SEARCH
750 =====
751 #!/usr/bin/env python3
752 # New file content here
753 def new_function():
754     pass
755 >>>>> REPLACE
756
757 Now, analyze the file content and instruction provided, and output
758 the necessary search-replace blocks.
759
760

```

A.3. System Prompt

Both baseline and SWE-Edit agents receive the same system prompt, which describes the task and provides step-by-step guidance. The template variables `{{ instance.repo_path }}` and `{{ instance.problem_statement }}` are populated for each SWE-bench instance.

SWE-Bench System Prompt

```

761 <uploaded_files>
762 {{ instance.repo_path }}
763 </uploaded_files>
764 I've uploaded a python code repository in the directory
765 {{ instance.repo_path }} (not in /tmp/inputs). Consider the
766 following issue descriptions:
767
768 <issue_description>
769

```

```

770
771     {{ instance.problem_statement }}
772   </issue_description>
773
774   Can you help me implement the necessary changes to the repository
775   so that the requirements specified in the <issue_description> are
776   met?
777   I've already taken care of all changes to any of the test files
778   described in the <issue_description>. This means you DON'T have to
779   modify the testing logic or any of the tests in any way!
780   Also the development Python environment is already set up for you
781   (i.e., all dependencies already installed), so you don't need to
782   install other packages.
783
784   Your task is to make the minimal changes to non-test files in the
785   {{ instance.repo_path }} directory to ensure the <issue_description>
786   is satisfied.
787
788   Follow these steps to resolve the issue:
789   1. As a first step, it might be a good idea to explore the repo to
790      familiarize yourself with its structure.
791   2. Create a script to reproduce the error and execute it with
792      'python <filename.py>' using the execute_bash tool to confirm
793      the error
794      - **Important:** If testing a Python package, add
795        'import sys; sys.path.insert(0, '{{ instance.repo_path }}')'
796        at the top of your script before package imports to ensure
797        you're testing the local version, not an installed version.
798   3. Edit the source code of the repo to resolve the issue
799   4. Rerun your reproduce script and confirm that the error is fixed!
800   5. Think about edge cases and make sure your fix handles them as
801      well
802
803   Your thinking should be thorough and so it's fine if it's very long.

```

B. Full Experiment Results

Table 7. Detailed Performance Metrics on SWE-bench Verified. Results are averaged over three independent runs for each configuration. “Succ.” denotes the success rate of editor tool calls.

Config.	Resolved (%)	Rounds	Agent Cost (\$)	Editor Cost (\$)	Viewer Cost (\$)	Total Cost (\$)	Output Tokens	Cached Input	Non-Cached Input	Viewer Calls	Editor Calls	Succ. (%)
Baseline	69.9	24.2	243.7	—	—	243.7	9632	369.8K	276.7K	5.78	2.86	93.4
+ Viewer	70.3	23.4	215.8	—	9.18	225.0	9708	312.9K	237.1K	4.26	2.75	94.3
+ Editor	71.3	22.7	263.0	5.28	—	268.3	8979	317.6K	318.2K	7.78	2.33	96.1
SWE-Edit	72.0	20.6	179.6	5.38	15.14	200.1	9517	304.6K	181.3K	7.49	2.37	96.6

C. PR-Edit Benchmark

This section provides implementation details for the PR-Edit Benchmark, including the normalization function used for computing the normalized match reward, the prompt used for GPT-4.1-based equivalence grading, and an example from the dataset.

C.1. Code Normalization

The normalized match reward compares model output against ground truth after canonicalizing whitespace and removing comments. This provides a reliable, execution-free proxy for edit correctness during training. Listing 3 shows the complete implementation.

```

825
826     Listing 3. Code normalization function for computing normalized match reward.
827
828     def normalize_code(code: str) -> str:
829         """
830             Normalize code by removing comments and normalizing whitespace.
831             This allows for comparison that tolerates comment and whitespace differences.
832
833             Note: This uses regex-based heuristics and may incorrectly handle
834             comment-like patterns inside string literals (e.g., "http://url" or "# not a comment").
835             For most code comparison tasks, this is an acceptable trade-off.
836         """
837
838         # Remove multi-line comments first (before single-line to handle edge cases properly)
839         # C-style /* */ comments
840         code = re.sub(r"/\*.*?\*/", "", code, flags=re.DOTALL)
841         # Python docstrings / multi-line strings used as comments
842         code = re.sub(r'""".*?"""', "", code, flags=re.DOTALL)
843         code = re.sub(r'\'\'\'.*?\'\'\'', "", code, flags=re.DOTALL)
844         # HTML/XML comments
845         code = re.sub(r"<!--.*?-->", "", code, flags=re.DOTALL)
846
847         # Remove single-line comments (// for C-like languages, # for Python/shell/etc.)
848         code = re.sub(r"//.*$", "", code, flags=re.MULTILINE)
849         code = re.sub(r"#.*$", "", code, flags=re.MULTILINE)
850
851         # Normalize all whitespace (spaces, tabs, newlines) to single space
852         # This collapses the code into a single line, ignoring all formatting differences
853         code = re.sub(r"\s+", " ", code)
854
855         # Strip leading/trailing whitespace
856         code = code.strip()
857
858     return code

```

C.2. GPT-4.1 Equivalence Grading

For the GPT Grader metric in Table 3, we prompt GPT-4.1 to assess whether the model’s edit is functionally equivalent to the ground truth. The grader receives the original code and two diffs (model output and ground truth), then determines logical equivalence while ignoring cosmetic differences such as formatting, comments, or variable naming.

GPT-4.1 Grader System Prompt

You are a code analysis expert specializing in logical equivalence comparison. You are given the original code and two diffs showing modifications to that same original code. Your task is to determine if these two modifications are functionally equivalent.

IMPORTANT:

- Focus on logical equivalence, not textual similarity
- Consider that different implementations can be functionally equivalent
- Ignore cosmetic differences like formatting, comments, or variable naming

Please analyze these diffs step by step, then provide your final answer.

REQUIRED OUTPUT FORMAT:

Analysis: [Your detailed analysis here]
Result: [EQUIVALENT/NOT_EQUIVALENT]

GPT-4.1 Grader User Prompt Template

Compare these two code modifications for logical equivalence:

ORIGINAL CODE:
{original_code}

DIFF 1:
{diff1}

DIFF 2:

```

880
881     {diff2}
882
883     Are these modifications functionally equivalent?
884
885

```

C.3. Dataset Example

Each instance in the PR-Edit Benchmark consists of three components: (1) the original file content before the pull request, (2) the ground truth file content after merging, and (3) a natural language edit query describing the required modification. Below is a representative example.

Edit Query

```

890
891
892     Replace the usage of 'assert_image_equal' with 'assert_image_equal_tofile',
893     and update imports accordingly to improve the way image comparisons are
894     being handled.
895

```

Original Code (excerpt)

```

896
897
898     import tempfile
899     from io import BytesIO
900
901     import pytest
902
903     from PIL import Image, ImageSequence, SpiderImagePlugin
904
905     from .helper import assert_image_equal, hopper, is_pypy
906
907     TEST_FILE = "Tests/images/hopper.spider"
908
909     ...
910
911     # for issue #4093
912     def test_odd_size():
913         data = BytesIO()
914         width = 100
915         im = Image.new("F", (width, 64))
916         im.save(data, format="SPIDER")
917
918         data.seek(0)
919         with Image.open(data) as im2:
920             assert_image_equal(im, im2)

```

Ground Truth (excerpt)

```

921
922
923     import tempfile
924     from io import BytesIO
925
926     import pytest
927
928     from PIL import Image, ImageSequence, SpiderImagePlugin
929
930     from .helper import assert_image_equal_tofile, hopper, is_pypy
931
932     TEST_FILE = "Tests/images/hopper.spider"
933
934     ...
935
936     # for issue #4093
937     def test_odd_size():

```

```

935     data = BytesIO()
936     width = 100
937     im = Image.new("F", (width, 64))
938     im.save(data, format="SPIDER")
939
940     data.seek(0)
941     with Image.open(data) as im2:
942         assert_image_equal_tofile(im, im2)

```

In this example, the edit requires two coordinated changes: updating the import statement and modifying the function call in `test_odd_size` (final line).

D. Open-Source Model Evaluation Details

D.1. Model Selection

Our main experiments use GPT-5, a proprietary model. To verify that SWE-Edit generalizes across model families, we evaluate on three recent open-source reasoning models: Kimi-K2-Thinking ([Moonshot AI, 2025](#)), MiniMax-M2.1 ([MiniMax AI, 2025](#)), and GLM-4.7 ([ZhipuAI, 2025](#)). These models were selected for two reasons: (1) they represent the latest generation of open-source models with strong reasoning capabilities, and (2) they have undergone substantial agentic training, making them suitable candidates for the challenging software engineering tasks.

D.2. Inference Configuration

All three models are configured with *Interleaved Thinking* and *Preserved Thinking* enabled. Interleaved Thinking allows the model to reason before every response and tool call, improving instruction following and generation quality. Preserved Thinking automatically retains reasoning blocks across multi-turn conversations, reusing existing reasoning rather than re-deriving from scratch—reducing information loss and improving consistency for long-horizon agentic tasks.

We use Fireworks AI² for model inference, following each model’s official inference settings on SWE-bench Verified. Common hyperparameters across all models include maximum new tokens of 16,384 and top- p of 1.0. Model-specific exceptions are noted in Table 8.

Table 8. Inference hyperparameters for open-source model evaluation. We follow official settings where available.

Model	Temperature	Top- p	Max Tokens
Kimi-K2-Thinking	1.0	1.0	16,384
MiniMax-M2.1	1.0	0.95	16,384
GLM-4.7	0.7	1.0	16,384

D.3. Evaluation Protocol

Due to computational constraints, we evaluate on the first 100 instances of SWE-bench Verified rather than the full 500-instance benchmark. Each configuration (baseline and SWE-Edit) is run twice to reduce variance. We use the same baseline agent scaffolding and SWE-Edit architecture as described in Appendix A, with only the main agent model swapped.

²<https://fireworks.ai>